

# Flink+TiDB

## 一、FlinkSQL+TiDB实现物化视图

在尝试各个版本的坑之后，最终选择Flink 1.13.6版本。组最简单的搭建方式。

TiDB集群中控机：172.16.11.70

```
1 [root@ogg bin]# hostname -I
2 172.16.11.70 172.17.0.1
3 [root@ogg bin]#
4 1.版本选择
5 [root@ogg bin]# ./flink -v
6 Version: 1.13.6, Commit ID: b2ca390
7 2.安装
8 tar xvf flink-1.13.6-bin-scala_2.11.tgz
9 cd flink-1.13.6/conf/
10 # 这里单机启动，ip等未做修改。
11 vim flink-conf.yaml
12 # 该参数如果是1，复杂任务无法完成
13 taskmanager.numberOfTaskSlots: 1 ---> 40
14 # 相应增加内存
15 jobmanager.memory.process.size: 16000m
16 taskmanager.memory.process.size: 17280m
17 3.在/home/ogg/flink-1.13.6/lib添加所需依赖的对应版本jar包
18 flink-connector-jdbc_2.11-1.13.6.jar
19 flink-sql-connector-tidb-cdc-2.2.0.jar
20 mysql-connector-java-8.0.30.jar
21 4.启动
22 [root@ogg bin]# pwd
23 /home/ogg/flink-1.13.6/bin
24 [root@ogg bin]# ./start-cluster.sh
25 5.起flink sql
26 [root@ogg bin]# pwd
27 /home/ogg/flink-1.13.6/bin
28 [root@ogg bin]# ./sql-client.sh
29 6.分别在tidb中创建test1, test2, users表
30 mysql> show create table test1 \G;
31 ***** 1. row *****
32      Table: test1
33 Create Table: CREATE TABLE `test1` (
34   `id` int(11) NOT NULL,
35   `name` varchar(50) DEFAULT NULL,
```

```

36     PRIMARY KEY (`id`) /*T![clustered_index] NONCLUSTERED */
37 ) ENGINE=InnoDB DEFAULT CHARSET=utf8mb4 COLLATE=utf8mb4_bin
38 1 row in set (0.00 sec)
39
40 mysql> show create table test2 \G;
41 ***** 1. row *****
42      Table: test2
43 Create Table: CREATE TABLE `test2` (
44   `id` int(11) NOT NULL,
45   `name` varchar(50) DEFAULT NULL,
46   PRIMARY KEY (`id`) /*T![clustered_index] NONCLUSTERED */
47 ) ENGINE=InnoDB DEFAULT CHARSET=utf8mb4 COLLATE=utf8mb4_bin
48 1 row in set (0.00 sec)
49
50 mysql> show create table users \G;
51 ***** 1. row *****
52      Table: users
53 Create Table: CREATE TABLE `users` (
54   `id` int(11) NOT NULL,
55   `name` varchar(50) DEFAULT NULL,
56   `id0` int(11) NOT NULL DEFAULT '1',
57   `name0` varchar(50) DEFAULT NULL,
58   PRIMARY KEY (`id`) /*T![clustered_index] CLUSTERED */
59 ) ENGINE=InnoDB DEFAULT CHARSET=utf8mb4 COLLATE=utf8mb4_bin
60 1 row in set (0.00 sec)
61
62 7. 分别在flink sql中创建test1, test2, users表对应的映射表test1, test2, t4
63 Flink SQL> SET 'execution.checkpointing.interval' = '1s';
64 Flink SQL> CREATE TABLE test1 (
65 >   id INT,
66 >   name STRING,
67 >   PRIMARY KEY(id) NOT ENFORCED
68 > ) with (
69 >   'connector' = 'tidb-cdc',
70 >   'tikv.grpc.timeout_in_ms' = '20000',
71 >   'pd-addresses' = '172.16.11.44:2381',
72 >   'database-name' = 'logadmin',
73 >   'table-name' = 'test1'
74 > );
75 Flink SQL> CREATE TABLE test2 (
76 >   id INT,
77 >   name STRING,
78 >   PRIMARY KEY(id) NOT ENFORCED
79 > ) with (
80 >   'connector' = 'tidb-cdc',
81 >   'tikv.grpc.timeout_in_ms' = '20000',
82 >   'pd-addresses' = '172.16.11.44:2381',

```

```

83 > 'database-name' = 'oggadmin',
84 > 'table-name' = 'test2'
85 > );
86
87 Flink SQL> CREATE TABLE t4 (
88 > id INT,
89 > name STRING,
90 > id0 INT,
91 > name0 STRING,
92 > PRIMARY KEY(id) NOT ENFORCED
93 > ) with (
94 > 'connector' = 'jdbc',
95 > 'url' = 'jdbc:mysql://172.16.11.43:4001/oggadmin',
96 > 'table-name' = 'users',
97 > 'driver' = 'com.mysql.cj.jdbc.Driver',
98 > 'username' = 'oggadmin',
99 > 'password' = 'oggadmin'
100 > );
101 8.将test1, test2表关联查询结果写入t4中, 同时写回到tidb的表users中。
102 Flink SQL> insert into t4 select * from test1,test2 where test1.id=test2.id;
103 [INFO] Submitting SQL update statement to the cluster...
104 [INFO] SQL update statement has been successfully submitted to the cluster:
105 Job ID: 8c7565ce11890f8f7584bf6ff2d47b7e
106 9.查询users中表数据
107 mysql> select * from users;
108 +-----+-----+-----+-----+
109 | id | name | id0 | name0 |
110 +-----+-----+-----+-----+
111 | 1 | 1 | 1 | 1 |
112 | 2 | 2 | 2 | 2 |
113 | 3 | 3 | 3 | 3 |
114 | 4 | 4 | 4 | 4 |
115 | 5 | 5 | 5 | 5 |
116 | 6 | 6 | 6 | 6 |
117 | 7 | 7 | 7 | 7 |
118 | 8 | 8 | 8 | 8 |
119 | 9 | 9 | 9 | 9 |
120 | 10 | 10 | 10 | 10 |
121 | 11 | 11 | 11 | 11 |
122 | 12 | 12 | 12 | 12 |
123 | 13 | 13 | 13 | 13 |
124 | 14 | 14 | 14 | 14 |
125 +-----+-----+-----+-----+
126 14 rows in set (0.00 sec)
127 mysql> select * from users;
128 +-----+-----+-----+-----+
129 | id | name | id0 | name0 |

```

```

130 +---+-----+-----+-----+
131 | 1 | 1      | 1 | 1      |
132 | 2 | 2      | 2 | 2      |
133 | 3 | 3      | 3 | 3      |
134 | 4 | 4      | 4 | 4      |
135 | 5 | 5      | 5 | 5      |
136 | 6 | 6      | 6 | 6      |
137 | 7 | 7      | 7 | 7      |
138 | 8 | 8      | 8 | 8      |
139 | 9 | 9      | 9 | 9      |
140 |10 |10      |10 |10      |
141 |11 |11      |11 |11      |
142 |12 |12      |12 |12      |
143 |13 |13      |13 |13      |
144 |14 |14      |14 |14      |
145 |15 |15      |15 |15      |
146 +---+-----+-----+-----+
147 15 rows in set (0.01 sec)

```

← → ↻ 不安全 | 172.16.11.70:8081/#/job/8c7565ce11890f8f7584bf6ff2d47b7e/overview

新标签页 PingCAP 数据库 DP 大数据 Gmail

Apache Flink Dashboard

Overview

Jobs

- Running Jobs
- Completed Jobs

Task Managers

Job Manager

Submit New Job

Version: 1.13.6 | Commit: b2ca390 @ 2022-02-03T14:54:22+01:00 | Message: 0

insert-into\_default\_catalog.default\_database.t4 **RUNNING** 5

ID: 8c7565ce11890f8f7584bf6ff2d47b7e | Start Time: 2022-09-01 14:07:51 | Duration: 17m 19s

Cancel Job

Overview Exceptions TimeLine Checkpoints Configuration

```

graph LR
    S1["Source: TableSourceScan(table=[default_catalog.default_database.test1], fields=[id, name])  
Parallelism: 1  
Backpressured (max): 0%  
Busy (max): N/A"]
    S2["Source: TableSourceScan(table=[default_catalog.default_database.test2], fields=[id, name])  
Parallelism: 1  
Backpressured (max): 0%  
Busy (max): N/A"]
    CN1["ChangelogNormalize(key=[id])  
Parallelism: 1  
Backpressured (max): 0%  
Busy (max): 0%"]
    CN2["ChangelogNormalize(key=[id])  
Parallelism: 1  
Backpressured (max): 0%  
Busy (max): 0%"]
    J["Join(joinType=[InnerJoin], where=[id = id0], select=[id, name, id0, name0], leftInputSpec=[JoinKeyContainsUniqueKey], rightInputSpec=[JoinKeyContainsUniqueKey]) -> NotNullEnforcer(fields=[id]) -> Sink: Sink(table=[default_catalog.default_database.t4], fields=[id, name, id0, name0])  
Parallelism: 1  
Backpressured (max): 0%  
Busy (max): 0%"]

    S1 -- HASH --> CN1
    S2 -- HASH --> CN2
    CN1 -- HASH --> J
    CN2 -- HASH --> J

```

```
[root@ogg lib]# ll
total 284548
-rwxrwxrwx 1 ogg oinstall      249570 Sep  1 13:05 flink-connector-jdbc_2.11-1.13.6.jar
-rwxrwxrwx 1 ogg oinstall       92314 Feb  4 2022 flink-csv-1.13.6.jar
-rwxrwxrwx 1 ogg oinstall 115425612 Feb  4 2022 flink-dist_2.11-1.13.6.jar
-rwxrwxrwx 1 ogg oinstall    148127 Feb  4 2022 flink-json-1.13.6.jar
-rwxrwxrwx 1 ogg oinstall   7709740 May  7 2021 flink-shaded-zookeeper-3.4.14.jar
-rwxrwxrwx 1 ogg oinstall  85355380 Aug 31 17:44 flink-sql-connector-tidb-cdc-2.2.0.jar
-rwxrwxrwx 1 ogg oinstall  36455408 Feb  4 2022 flink-table_2.11-1.13.6.jar
-rwxrwxrwx 1 ogg oinstall  41077430 Feb  4 2022 flink-table-blink_2.11-1.13.6.jar
-rwxrwxrwx 1 ogg oinstall    208006 Jan 13 2022 log4j-1.2-api-2.17.1.jar
-rwxrwxrwx 1 ogg oinstall    301872 Jan  7 2022 log4j-api-2.17.1.jar
-rwxrwxrwx 1 ogg oinstall   1790452 Jan  7 2022 log4j-core-2.17.1.jar
-rwxrwxrwx 1 ogg oinstall    24279 Jan  7 2022 log4j-slf4j-impl-2.17.1.jar
-rwxrwxrwx 1 ogg oinstall   2513563 Sep  1 11:16 mysql-connector-java-8.0.30.jar
[root@ogg lib]# pwd
/home/ogg/flink-1.13.6/lib
[root@ogg lib]#
```